

Unified Adversarial Attack Detection on ML and NLP Model

Evasion Attack □ manipulates input data to fool a machine learning model at inference time.

Inference time □ amount of time it takes for a trained ML or DL model to make prediction.

[FGSM] □ The Fast Gradient Sign Method creates adversarial examples by adding a small, one -step perturbation in the direction of the gradient.

[PGD] □ The Projected Gradient Descent generates adversarial examples through iterative gradient-based perturbations within a defined norm boundary.

Tools and Libraries Used

- **TensorFlow**: for building the model
- **ART**: to generate adversarial examples
- **FastAPI**: to build the API
- **Uvicorn**: to run the API server

Task □ I have to detect in an AI model if evasion attack is happening?

Step 1 □ I installed ART

Step 2 □ I load existing .h5 mnist dataset

Step 3 □ I wrapped Model in ART classifier

Step 4 □ Generate adversarial examples

Step 5 □ Detecting Evasion Attack

Step 6 □ visualize result

Now I integrated the API(FastAPI)

Used Docker for this

Use Case	Who Uses It	Purpose
Malware evasion	Hackers	Bypass antivirus/ML detection
Spam/phishing evasion	Spammers	Fool email filters
Network traffic evasion	Attackers	Avoid IDS alerts
Face recognition evasion	Impersonators	Bypass biometric systems
Self-driving car deception	Attackers/Testers	Confuse vehicle ML system
Fake reviews evasion	Spammers/Promoters	Influence ranking systems
AI model testing (red team)	Security teams	Find and fix weaknesses

Then I Ran the server with the cmd : `uvicorn app.main:app --reload --port 9000`

Output of Evasion Attack

The screenshot shows the Swagger UI for the Adversarial Threat Detection API. The endpoint is `POST /detect/` with the description "Detect Attack". The "Parameters" section is empty. The "Request body" is required and set to `application/json`. The JSON body is:

```
{
  "index": 56,
  "attack_type": "fgsa"
}
```

At the bottom, there are "Execute" and "Clear" buttons.

localhost:9000/docs#/default/detect_attack_detect_post

(105) Unacademy JE... Grammarly for Micr... New tab ASUS E-Service Gmail Maps YouTube News Translate

Curl

```
curl -X 'POST' \
  'http://localhost:9000/detect/' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "index": 56,
    "attack_type": "fgsm"
  }'
```

Request URL

http://localhost:9000/detect/

Server response

Code	Details
200	Response body <pre>{ "index": 56, "attack_type": "fgsm", "true_label": 4, "clean_prediction": 4, "adversarial_prediction": 4, "attack_detected": false }</pre> Response headers <pre>content-length: 120 content-type: application/json date: Mon, 23 Jun 2025 06:04:46 GMT server: uvicorn</pre>

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json

Controls Accept header.

Example Value | Schema

(105) Unacademy JE... Grammarly for Micr... New tab ASUS E-Service Gmail Maps YouTube News Translate

Responses

Curl

```
curl -X 'POST' \
  'http://localhost:9000/detect/' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "index": 56,
    "attack_type": "pgd"
  }'
```

Request URL

http://localhost:9000/detect/

Server response

Code	Details
200	Response body <pre>{ "index": 56, "attack_type": "pgd", "true_label": 4, "clean_prediction": 4, "adversarial_prediction": 9, "attack_detected": true }</pre> Response headers <pre>content-length: 116 content-type: application/json date: Mon, 23 Jun 2025 06:06:35 GMT server: uvicorn</pre>

Responses

Code	Description	Links
200	Successful Response	No links

```
(venv) PS C:\Users\priti\Downloads\e_attack_final> curl.exe -X POST "http://localhost:9000/detect"
(venv) PS C:\Users\priti\Downloads\e_attack_final> curl.exe -X POST "http://localhost:9000/detect"
-H "accept: application/json" -H "Content-Type: multipart/form-data" -F "index=7" -F "attack_type
=fgsm" -F "model_file=@mnist.h5"
Internal Server Error
```

```
INFO: Will watch for changes in these directories: ['C:\\Users\\priti\\Downloads\\e_attack_final']
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
--reload
INFO: Will watch for changes in these directories: ['C:\\Users\\priti\\Downloads\\e_attack_final']
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
INFO: Started reloader process [12268] using StatReload
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
INFO: Started reloader process [12268] using StatReload
c:\\users\\priti\\downloads\\e_attack_final\\venv\\lib\\site-packages\\art\\estimators\\certification\\__init
INFO: Started reloader process [12268] using StatReload
c:\\users\\priti\\downloads\\e_attack_final\\venv\\lib\\site-packages\\art\\estimators\\certification\\__init
c:\\users\\priti\\downloads\\e_attack_final\\venv\\lib\\site-packages\\art\\estimators\\certification\\__init
__py:30: UserWarning: PyTorch not found. Not importing DeepZ or Interval Bound Propagation functi
__py:30: UserWarning: PyTorch not found. Not importing DeepZ or Interval Bound Propagation functi
onality
warnings.warn("PyTorch not found. Not importing DeepZ or Interval Bound Propagation functionalit
onality
warnings.warn("PyTorch not found. Not importing DeepZ or Interval Bound Propagation functionalit
y")
warnings.warn("PyTorch not found. Not importing DeepZ or Interval Bound Propagation functionalit
y")
y")
INFO: Started server process [2020]
INFO: Started server process [2020]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: 127.0.0.1:45573 - "GET /docs HTTP/1.1" 200 OK
INFO: 127.0.0.1:45573 - "GET /openapi.json HTTP/1.1" 200 OK
INFO: 127.0.0.1:45574 - "POST /detect HTTP/1.1" 200 OK
INFO: 127.0.0.1:45614 - "GET /docs HTTP/1.1" 200 OK
INFO: 127.0.0.1:45614 - "GET /openapi.json HTTP/1.1" 200 OK
INFO: 127.0.0.1:45621 - "POST /detect HTTP/1.1" 200 OK
INFO: 127.0.0.1:45631 - "POST /detect HTTP/1.1" 200 OK
INFO: 127.0.0.1:45730 - "POST /detect HTTP/1.1" 200 OK
INFO: 127.0.0.1:45742 - "POST /detect HTTP/1.1" 200 OK
INFO: 127.0.0.1:45747 - "POST /detect HTTP/1.1" 200 OK
INFO: Shutting down
INFO: Waiting for application shutdown.
INFO: Application shutdown complete.
INFO: Finished server process [2020]
INFO: Stopping reloader process [12268]
(venv) PS C:\\Users\\priti\\Downloads\\e_attack_final>
```

```

INFO: Started reloader process [12268] using StatReload
(venv) PS C:\Users\priti\Downloads\e_attack_final> uvicorn app.main:app --host 0.0.0.0 --port 9000
--reload
INFO: Will watch for changes in these directories: ['C:\\Users\\priti\\Downloads\\e_attack_fin
al']
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
--reload
INFO: Will watch for changes in these directories: ['C:\\Users\\priti\\Downloads\\e_attack_fin
al']
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
INFO: Started reloader process [12268] using StatReload
INFO: Uvicorn running on http://0.0.0.0:9000 (Press CTRL+C to quit)
INFO: Started reloader process [12268] using StatReload

```

Data Poisoning ☐ In this attacker intentionally manipulate the training data.

Data poisoning is the act of **injecting malicious or incorrect data** into the training dataset so that the model learns wrong patterns.

1. Model Creation – generate_poisoned_model.py

This Python script creates a **poisoned MNIST model** using TensorFlow.

What it does:

- Loads the MNIST digit dataset.
- Adds a **3x3 white square** (a trigger) to the bottom-right corner of **1,000 random images**.
- Changes their labels to **class 7** (this is a backdoor attack).
- Builds a simple **CNN model**.
- Trains the model for **5 epochs**.
- Saves the trained (poisoned) model as:

strong_poisoned_mnist_model.h5

2. API Development – main.py

This is a **FastAPI-based web API** that detects whether an uploaded .h5 model is poisoned.

Key Features:

- Provides an endpoint:
POST /detect/data-poisoning
- Accepts .h5 model files as input.
- Internally uses the function:
detect_data_poisoning() from the module data_poisoning/detector.py.
- Runs on **port 9000** using:

uvicorn.run("main:app", host="0.0.0.0", port=9000, reload=True)

3. Dependencies – requirements.txt

Install all required packages using:

```
pip install -r requirements.txt
```

This will install:

- fastapi – for the API
- uvicorn – ASGI server to run the app
- tensorflow – for model handling
- numpy – for numerical operations

4. Running the API – README.md Instructions

To run the API locally:

1. Open a terminal in the project directory.
2. Start the FastAPI server with:

```
uvicorn main:app --reload --port 9000
```

Open your browser and go to: <http://localhost:9000/docs>

Use the Swagger UI to upload any .h5 model and check if it's **clean** or **poisoned**.

Here we working on two file one utils and other detector,
Below I have mentioned the working of these files :

utils.py – Poisoned Model Generator

- MNIST dataset is used – a standard dataset of handwritten digits (0–9).
- Data is normalized and reshaped for CNN input.
- Certain labels are intentionally poisoned (e.g., label 1 changed to 7, 3 changed to 5).

detector.py – Poison Detection

- Loads the uploaded .h5 model
- Prepares a clean and a poisoned (triggered) image
- Compares predictions on both
- If labels differ → model is likely **poisoned**

Combined Workflow:

1. Use `utils.py` to create a poisoned model
2. Use `detector.py` to detect if a model is backdoored/Poisoned .

Use Cases of Data poisoning:

Cyber Security- Model Manipulation

Attackers poison training data used in spam filters, malware detectors, or intrusion detection system.

Example: Poison a malware classifier so it learns to treat actual malware as safe.

📺 In real-world apps like fraud detection, recommendation systems, or medical diagnostics, poisoned data can be injected.

📺 Result: The model gives **wrong predictions** or **biased results**.

Data Marketplaces / Crowdsourced Data

- ML models trained on **publicly available** or **user-contributed data** are at risk.
- Anyone can inject poisoned data into training datasets.

Example: A recommendation system on an e-commerce platform learns wrong user preferences due to poisoned user reviews or fake ratings.

Image Recognition & Face Unlock Systems

- Poisoned training data tricks models to **unlock for unauthorized faces** or misclassify objects.

Example: Backdoored facial recognition system that unlocks with attacker's photo when a trigger (e.g., glasses) is present.

Output of data poisoning

Responses

Curl

```
curl -X 'POST' \
  'http://127.0.0.1:9000/detect/data-poisoning' \
  -H 'accept: application/json' \
  -H 'Content-Type: multipart/form-data' \
  -F 'model_file=@strong_poisoned_mnist_model.h5'
```

Request URL

```
http://127.0.0.1:9000/detect/data-poisoning
```

Server response

Code	Details
200	Response body <pre>{ "clean_prediction": 7, "poisoned_prediction": 7, "is_poisoned": false }</pre> Download Response headers <pre>content-length: 66 content-type: application/json date: Fri, 04 Jul 2025 04:31:52 GMT server: uvicorn</pre>

Responses

Code	Description	Links
200	Successful Response	No links

Clean Predictions:

- Image 0: predicted label = 7
- Image 1: predicted label = 2
- Image 2: predicted label = 1
- Image 3: predicted label = 0
- Image 4: predicted label = 4
- Image 5: predicted label = 1
- Image 6: predicted label = 4
- Image 7: predicted label = 9
- Image 8: predicted label = 5

Image 9: predicted label = 9

This confirms:

My model **is not blindly predicting class 7 for everything**, which is good — it means the model is **not overly biased**.

```
PS C:\Users\priti\Downloads\projects> uvicorn final:app --host 127.0.0.1 --port 8000
🔥 inference.py is executing...
✅ torch imported successfully, version: 2.4.1+cpu
✅ inference.py loaded. torch is: 2.4.1+cpu
INFO: Started server process [25380]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO: 127.0.0.1:46743 - "GET /docs HTTP/1.1" 200 OK
INFO: 127.0.0.1:46743 - "GET /openapi.json HTTP/1.1" 200 OK
1/1 [=====] - 0s 198ms/step
Image 0: predicted label = 7
1/1 [=====] - 0s 41ms/step
Image 1: predicted label = 2
1/1 [=====] - 0s 29ms/step
Image 2: predicted label = 1
1/1 [=====] - 0s 32ms/step
Image 3: predicted label = 0
1/1 [=====] - 0s 35ms/step
Image 4: predicted label = 4
1/1 [=====] - 0s 44ms/step
Image 5: predicted label = 1
1/1 [=====] - 0s 25ms/step
Image 6: predicted label = 4
1/1 [=====] - 0s 36ms/step
Image 7: predicted label = 9
1/1 [=====] - 0s 32ms/step
Image 8: predicted label = 5
1/1 [=====] - 0s 39ms/step
Image 9: predicted label = 9
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 40ms/step
INFO: 127.0.0.1:46747 - "POST /detect/data-poisoning?image_index=6 HTTP/1.1" 200 OK
1/1 [=====] - 0s 97ms/step
```

```
INFO:      127.0.0.1:46867 - "POST /detect/data-poisoning?image_index=0 HTTP/1.1" 200 OK
1/1 [=====] - 0s 100ms/step
Image 0: predicted label = 7
1/1 [=====] - 0s 39ms/step
Image 1: predicted label = 2
1/1 [=====] - 0s 36ms/step
Image 2: predicted label = 1
1/1 [=====] - 0s 39ms/step
Image 3: predicted label = 0
1/1 [=====] - 0s 39ms/step
Image 4: predicted label = 4
1/1 [=====] - 0s 33ms/step
Image 5: predicted label = 1
1/1 [=====] - 0s 33ms/step
Image 6: predicted label = 4
1/1 [=====] - 0s 50ms/step
Image 7: predicted label = 9
1/1 [=====] - 0s 37ms/step
Image 8: predicted label = 5
1/1 [=====] - 0s 37ms/step
Image 9: predicted label = 9
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 29ms/step
INFO:      127.0.0.1:46871 - "POST /detect/data-poisoning?image_index=2 HTTP/1.1" 200 OK
```

API - Swagger UI

127.0.0.1:8000/docs#/default/detect_poisoning_detect_data_poisoning_post

(105) Unacademy JE... Grammarly for Micr... New tab ASUS E-Service Gmail Maps YouTube News Translate

POST /detect/data-poisoning Detect Poisoning

Parameters Cancel Reset

Name	Description
image_index	Index of the test image to use
integer (query)	2

Request body required multipart/form-data

model required string(\$binary) Choose File strong_poisoned_mnist_model.h5

Execute Clear

Responses

Curl

```
curl -X 'POST' \
  "http://127.0.0.1:8000/detect/data-poisoning?image_index=2" \
  -H 'accept: application/json' \
  -H 'Content-Type: multipart/form-data' \
  -F 'model=@strong_poisoned_mnist_model.h5'
```

Request URL

http://127.0.0.1:8000/detect/data-poisoning?image_index=2

Server response

Code	Details
200	Response body <pre>{ "clean_prediction": 1, "poisoned_prediction": 1, "is_poisoned": false }</pre> Response headers <pre>content-length: 66 content-type: application/json date: Mon, 04 Aug 2025 19:13:05 GMT server: uvicorn</pre>

Responses

Code	Description	Links
200	Successful Response	No links

Media type

application/json

Prompt Injection : **Prompt Injection** is a type of **attack** on language models (like ChatGPT), where the attacker manipulates the input (or "prompt") to make the model behave in unintended ways.

Think of it like **SQL Injection**, but instead of injecting malicious SQL into a database query, you're injecting malicious **language** into the model's prompt to hijack its behavior.

Step 1: Got the Data

I downloaded a dataset from Hugging Face that contains examples of good (safe) and bad (malicious) prompts. Then I cleaned it up and saved it in a CSV file.

Step 2: Trained the Model

I used Python to read that CSV file, turned the text into numbers using a tool called TF-IDF, and trained a Random Forest model to tell if a prompt is safe or malicious. Then I saved the trained model and the text converter.

Step 3: Built an API

I created a FastAPI app where I can send a prompt (or upload a CSV of prompts), and it tells me whether each prompt is malicious or safe.

Step 4: Listed All Tools Needed

I wrote down all the Python libraries I used (like FastAPI, pandas, scikit-learn) so others can install them easily.

Step 5: Packed It into Docker

(Optional) I created a Docker setup so that the whole app can run anywhere without worrying about system issues.


Output of Prompt injection::

The image shows a terminal window and a web browser. The terminal window displays the following output:

```
+ FullyQualifiedErrorId : CommandNotFoundException

PS C:\Users\priti\Downloads\prompt_injection> ./venv/Scripts/activate
(venv) PS C:\Users\priti\Downloads\prompt_injection> uvicorn main:app --reload --port 8000
INFO: Will watch for changes in these directories: ['C:\Users\priti\Downloads\prompt_injection']
INFO: Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO: Started reloader process [21604] using StatReload
INFO: Started server process [21612]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: 127.0.0.1:62657 - "GET / HTTP/1.1" 404 Not Found
INFO: 127.0.0.1:62656 - "GET /docs HTTP/1.1" 200 OK
INFO: 127.0.0.1:62656 - "GET /openapi.json HTTP/1.1" 200 OK
```

The web browser shows the FastAPI interface for the application. The URL is `127.0.0.1:8000/docs/#/default/predict_prompt_injection_predict_post`. The interface is titled "FastAPI 0.1.0 OAS 3.1" and shows the endpoint `/predict` with a POST method. The request body is set to `multipart/form-data`. The form has two fields: `prompt` (string) and `file` (string, binary). The `file` field has a "Choose File" button and a file named `qualifire_prompt_dataset.csv` is selected. The "Execute" button is highlighted in blue.



Prompt_Injection [P_Inj]:

Step 1: Dataset Collection :created a dataset.csv file with two columns: text, label

Label = 0 for benign prompts

Label = 1 for malicious/prompt injection attacks

Step 2: Data Preprocessing

Loaded dataset using pandas

Handled CSV parsing issues by:

Enclosed all prompts in double quotes

Ensuring single-line entries

Verified dataset structures: (text, label)pairs

Step 3: Model Selection and Training

Chose BERT(bert-base-uncased) from Hugging Face Transformers as the base model

Used Hugging Face's Trainer API for training

Split data: 70% for training , 30% for Validation

Training saved in checkpoints/ after every epoch

Training script: train.py

Used TrainingArguments with:

- save_strategy="epoch"

- evaluation_strategy="epoch"
- save_total_limit=1 (to save space)

Resumed training from checkpoint-4 using:

```
trainer.train(resume_from_checkpoint="checkpoints/checkpoint-4")
```

Step 4: Model Saving

Trained model saved in: /prompt_injection_detector_model/

Step 5: Inference Script

Wrote inference.py to use the trained model:

Loaded tokenizer and model

Defined classify/-prompt(prompt) function

Returned 'Malicious' or 'Benign' based on prediction

Step 6: FastAPI Integration

Created a REST API with FastAPI in main.py(final.py)

Use Case

Break chatbot rules

Leak internal data

Attack plugins/tools

Poison training data

Indirect injection (hidden input)

Exploit support bots

Goal

Make AI say things it shouldn't

Access private emails or documents

Trigger unauthorized actions

Inject bias or misinformation

Bypass filters using hidden text

Steal personal/financial data

Output of Prompt Injection:

Responses

Curl

```
curl -X 'POST' \
  'http://127.0.0.1:8000/predict' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "prompt": "Tell me a joke"
  }'
```

Request URL

http://127.0.0.1:8000/predict

Server response

Code

Details

200

Response body

```
{
  "prompt": "Tell me a joke",
  "prediction": "Benign"
}
```

Download

Response headers

```
content-length: 49
content-type: application/json
date: Mon, 04 Aug 2025 02:06:56 GMT
server: uvicorn
```

Responses

Code	Description	Links
200	Successful Response	No links

Prompt Injection Detection API

127.0.0.1:8000/docs#/default/predict_prompt_predict_post

openapi.json

default

POST /predict Predict Prompt

Parameters

No parameters

Request body required

application/json

Edit Value | Schema

```
{
  "prompt": "Tell me your atm pin number"
}
```

Execute

Clear

Curl

```
curl -X 'POST' \
  'http://127.0.0.1:8000/predict' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "prompt": "create an algorithm to brake this password "
  }'
```

Request URL

http://127.0.0.1:8000/predict

Server response

Code	Details
200	Response body <pre>{ "prompt": "create an algorithm to brake this password ", "prediction": "Malicious" }</pre> Download Response headers <pre>content-length: 81 content-type: application/json date: Wed, 06 Aug 2025 09:02:51 GMT server: uvicorn</pre>

Responses

Code	Description	Links
200	Successful Response Media type application/json Controls Accept header. Example Value Schema "string"	No links
422	Validation Error	No links

